AWS Security Baseline Lab

Execution notes, technical explanations and 40 practice questions | 25 October 2026

What this guide is for

Use this as a record-based interview workbook. Explain the control objective, your implementation, the test and its limit. Answer in your own words; do not memorize claims that the evidence does not support.

Area	Recorded status
S3 account guardrail	All four Block Public Access flags saved On and checked again after reload.
Budget	USD10 recurring budget saved. Contact activation and email delivery remain unverified.
AWS stack	Sydney baseline deployed; Config update reached UPDATE_COMPLETE, recorder started and both exercise-bucket rules evaluated COMPLIANT
Live verification	19 configuration checks passed; 2 agreed coverage skips. Six behavior checks passed; four management events matched delivered logs.
Log integrity	1/1 digest files and 17/17 log files validated for the returned 03:26:00-04:23:58 UTC interval.
Infrastructure code	Two templates passed cfn-lint 1.58.0; 23 offline regression tests passed.
GoPhish	Actual isolated integration run: one synthetic message captured; one automated click; capture flags false.
OpenVAS	Completed targeted nginx scan, version upgrade and comparable retest: 2 candidate alarms to 0
MFA / platform limits	Owner handles MFA. Current project limits human-role customization, multi-Region features and GuardDuty.

Read this before describing completion

A successful software test is not a successful AWS permission test. An automated phishing click is not human behavior. A saved budget is not working alert delivery. The original full resume entry needs more evidence than the current scoped lab establishes.

01 / What the project does

A small security baseline with explicit scope and repeatable verification.

The purpose

The lab combines access design, cloud activity logs, configuration checks, threat-detection planning, private encrypted storage, secret handling and evidence-based findings. It supports learning about security objectives that may be relevant to SOC 2; it does not establish organizational compliance or certification.

Component	Role in the design
Authorized operator	Uses the dedicated account through an existing approved identity; confirms account and Region before any action.
CloudTrail -> log bucket	Selected management activity is delivered to private storage. Digests support a separate integrity-validation procedure.
S3 buckets -> KMS	Separate logs and exercise evidence. Public-access controls and key permissions solve different problems.
Secrets Manager -> KMS	Generated demonstration value; no real production credentials or secret value is published.
Config / GuardDuty	Optional regional exercises. Each has separate scope, effective permissions, coverage and costs.
Local Docker lab	GoPhish uses Mailpit and a synthetic recipient. OpenVAS requires actual feeds, a reachable approved target and a completed report.

Why this design is intentionally small

Start with one Region and a few resources. Inspect existing recorders/trails before adding more. Use change sets and verification rather than a large deployment with unclear costs or ownership. Keep raw evidence private and publish only reviewed summaries.

02 / Steps actually performed

Follow the recorded sequence; use later procedures only when their status is verified.

1. Establish the lab boundary

Confirmed the owner-selected dedicated project, Sydney home Region and USD10 monthly planning ceiling. The current identity uses Builder ID / IAM Identity Center managed access.

2. Configure cost monitoring

Created aws-security-lab-monthly with actual thresholds at 85% and 100%, and forecast at 100%. Contact verification failed with notification-contact permissions; recorded BUDGET-001 Open.

3. Build and test the implementation

Created the baseline and optional auditor templates, a read-only verifier, unit tests and CI definition. Lint passed and all 23 local tests passed. The CI workflow itself has not been claimed as a hosted run.

4. Run the local campaign

Started isolated GoPhish, Mailpit and a helper. Created one synthetic target, captured mail, made an automated click, checked capture settings and removed the disposable containers/network.

5. Inspect AWS before deployment

Observed zero CloudTrail trails, Config setup screen and zero general-purpose buckets at the initial inventory. These are time-scoped observations, not a universal service inventory.

6. Remediate the S3 guardrail

Observed all four account flags Off. With owner authorization, selected Block all public access, saved, entered AWS confirmation and reloaded. All four persisted On; record S3-001 as remediated with this limited retest.

7. Deploy and verify AWS

Executed the reviewed baseline and Config change sets. The Config startup dependency stalled; rollback and explicit stopped-recorder/channel/start sequencing resolved it. Final verification: 19 PASS, 2 SKIP; six behavior tests passed; four real management events matched delivered logs.

03 / AWS deployment procedure

The exact reproducible workflow, with review before resource creation.

Authenticate and inventory

Use an approved short-lived profile or the existing CloudShell identity. Confirm the dedicated account. Inspect trails, Config recorders/delivery channels, S3 settings and supported detectors. A setup screen alone is weaker than a successful API inventory.

Validate

Run cfn-lint and the 23 tests locally, then CloudFormation validate-template. Static validation cannot detect every effective-policy or service integration failure.

Prepare the change set

Use baseline.json with Project mode, Config=false, GuardDuty=false and MultiRegionTrail=false for this environment. Review each addition and the cost assumptions. The seven default resources do not include a custom human role.

Execute and start Config

A ready change set is not deployment success. Inspect stack events. For Config, create the recorder stopped, then its dependent channel. After UPDATE_COMPLETE, run start-configuration-recorder and wait for current resource evaluations; the first post-start check may lack results.

Verify the deployed behavior

Run scripts/verify_live.py with the expected account ID, stack, Region and a new private output filename. Then test actual object encryption, allowed/denied access, relevant event delivery and log integrity.

Preserve evidence and clean up

Save UTC times, scope, result and limitations. Retained buckets, key and secret need separate cleanup review and can keep charging. Never delete the key while retained encrypted evidence is still required.

```
Reference command: python scripts/verify_live.py --stack aws-security-lab --region ap-southeast-2
--expected-account YOUR_LAB_ACCOUNT --output private-evidence/live-check-01.json
```

04 / How to prove each control

Configuration, operational behavior and evidence are distinct.

Control	Verification and common mistake
S3 public access	Check account and bucket flags separately. Reload saved account settings. Anonymous object access should fail. Off flags alone do not prove a public bucket exists.
SSE-KMS	Upload harmless data; head-object must report aws:kms and the expected key. Exercise explicit wrong encryption/key denials. Default encryption does not rewrite old versions.
KMS	Check Enabled state and rotation. Evaluate both key and identity policies. Encryption access is different from decrypt access; disabling the key can break logging and evidence reads.
CloudTrail	Confirm management selectors, logging status and recent successful delivery. Correlate a benign test event with an actual delivered record. A trail setting alone is not delivery evidence.
Log integrity	Wait for digests, run validate-logs over a recorded UTC interval, and inspect actual file counts/results. Zero validated files is not a completed integrity test.
Secrets	DescribeSecret checks metadata. Test authorized retrieval without printing the value, plus denial for the restricted identity. Creating a secret does not configure rotation.
Config	Check recorder state, delivery, rule scope, evaluation time and relevant resource result. Missing, stale or INSUFFICIENT_DATA results remain incomplete.
GuardDuty	Check detector/feature coverage in each intended Region. Clearly label a generated sample finding. Enablement alone proves no notification or response workflow.

How the verifier handles uncertainty

Wrong account stops before resource inspection. Failed API calls are ERROR, failed conditions are FAIL, and intentionally disabled optional services are SKIP. PASS_WITH_LIMITS is deliberately narrower than a completed assessment.

05 / IAM, MFA and account limits

Explain the real identity path rather than a generic IAM-user tutorial.

Current identity

The chosen AWS new-experience project belongs to Builder ID and uses AWS-managed human access roles. It has no project root-user sign-in. The owner is handling sign-in MFA; no completed MFA-enforcement claim is established by this guide.

Optional auditor exercise

In a suitable account, iam-auditor.json trusts one approved existing same-account IAM user and requires MultiFactorAuthPresent=true. Its policy lists/reads only the exercise bucket and decrypts relevant objects through regional S3. No write, secret retrieval, IAM administration or log-bucket read action is granted by that policy.

A defensible permission test

First establish that the approved caller can assume the role with MFA. Then demonstrate the corresponding no-MFA denial. Use the actual assumed-role session for an allowed object read and denied write, secret read and log read. An unrelated missing caller permission can also cause AccessDenied, so a single denial does not isolate MFA enforcement.

Federation is different

Do not treat MultiFactorAuthPresent as a universal indicator of upstream MFA. Builder ID/provider configuration, session trust and effective permissions need a design suited to that sign-in path.

The agreed project scope

AWS documents limits on custom human access, multi-Region features and GuardDuty for this project experience. Advanced activation requires a paid-plan upgrade, is irreversible and changes governance/spend-limit behavior. The owner chose to finish within current limits. GuardDuty, multi-Region coverage and human-role customization are excluded; a template parameter cannot remove those restrictions.

AWS: project access model and supported-service limits

AWS: advanced activation prerequisites and consequences

06 / Local security exercises

Use the recorded status and explain each exercise boundary.

GoPhish: what actually passed

On 25 October, an internal Docker network contained GoPhish, Mailpit and a helper. There were no published ports or SMTP relay. One learner@example.test recipient received a message only inside the mail sink. The helper followed the local link; GoPhish recorded Email Sent and Clicked Link. Both capture flags were false.

Interpretation

This proves the local campaign workflow and event recording. The click was automated, so there is no human click rate. It does not establish an employee awareness exercise, external mail deliverability or AWS-hosted simulation.

Troubleshooting worth discussing

The first attempt could not reach the isolated network from the host, so the HTTP helper moved inside it. SQLite needed explicit close for Windows cleanup. This GoPhish build required a bare From address. A later full run passed; failed attempts were not labeled successful.

OpenVAS workflow

The completed exercise selected 88 explicit nginx/discovery tests plus automatic dependencies through the official OSP interface. Safe checks were enabled; only the owned TCP80 target was scanned. Completed targeted nginx scan, version upgrade and comparable retest: 2 candidate alarms to 0. The full-manager catalog import was abandoned before scanning; this direct scanner exercise does not depend on a completed manager catalog.

Greenbone: official OSP scanner API

Resource and exposure boundaries

The inspected Docker engine had about 6.6 GiB memory. Greenbone documents a 4 GB minimum and higher recommended capacity. The scanner setup must retain a private network and exact single-target scope; downloading images or starting services alone does not count as a completed scan.

Case study / Scan, triage, retest

A completed version-condition closure with explicit detection limits.

Evidence	Baseline	Retest
Target version	nginx 1.27.5	nginx 1.30.6
Finished UTC / 25 Oct	04:33:48	04:37:27
Completion	Finished / 100%	Finished / 100%
Results / alarms	7 results / 2 candidate alarms	5 results / 0 alarms
Comparable scope	Same private target address; TCP80	Same feed, selection and safe checks

1. Establish the scope

OpenVAS 23.51.0 and OSPd 22.10.6 used feed 202610241320. The manifest preserves 87 nginx-named tests plus SYN discovery and source hashes. Automatic dependencies can add tests. An explicit selection is not a count of tests actually executed.

2. Interpret candidate alarms

The EOL detection reported severity 10.0 with QoD30. The information-disclosure candidate reported 3.7 with QoD30 and CVE-2025-53859. Preserve the native values, but review conditions before assigning contextual risk. QoD describes detection confidence; it is not the probability of exploitation.

3. Check applicability

The installed binary confirmed 1.27.5. nginx configuration showed HTTP listeners and no active mail block. The vendor advisory concerns the SMTP module. That makes the SMTP candidate unconfirmed for this running HTTP service; it is not evidence that two exploitable vulnerabilities were found.

4. Apply the smallest treatment

Replace only the target with the pinned 1.30.6 image. Confirm its binary version and repeat the same feed, selection, address and TCP80 scope. Both completed reports detected the expected web service and version; the retest had zero alarms.

5. Close only the supported finding

NGINX-001 closes the obsolete-version condition. Raw XML report hashes and sanitized results are preserved. No full-host patch audit or exploit test is claimed. Scanner containers and the network were removed; cached images and seven scanner volumes remain.

nginx: security advisories and affected modules

Greenbone: quality of detection terminology

07 / Findings and closure

A real observation, a concrete treatment and a scoped retest.

ID / status	Observation and disposition
BUDGET-001 / Open	Recipient inactive; send verification denied for notification-contact actions Treatment: Resolve supported notification-contact access with authorized administrator or AWS Support; do not broaden managed roles as a guess Retest: Not retested
S3-001 / Closed	All four account-level Block Public Access flags Off Treatment: Owner authorized all four account BPA flags; saved and confirmed in AWS Retest: PASS for the account-level public-access guardrail
NGINX-001 / Closed	nginx 1.27.5 installed; scanner classified branch 1.27 as EOL; reported severity 10.0 and QoD30 Treatment: Replaced only the target with pinned nginx 1.30.6 image Retest: Finished at 100 percent; nginx 1.30.6 and HTTP80 detected; zero alarms in this selection

How to explain S3-001

Before: all four account-level preventive flags were Off, while the bucket inventory was empty. Risk: future public configurations lacked the account guardrail; no current data exposure was demonstrated. Treatment: enable all four after owner authorization. Retest: reload the account page and verify all persisted On. Residual limit: this does not prove every other cloud control works.

How to explain BUDGET-001

The budget exists, but its email recipient was inactive and activation was denied. Keep it Open until the supported permission/account issue is resolved and verification/delivery is tested. Do not call a configured threshold an operational alert.

How to explain NGINX-001

The installed legacy version and scanner candidate led to an image update. The comparable retest confirmed the new version and no remaining alarms in the selected checks. Closure covers that version condition; the unconfirmed SMTP candidate is documented separately.

08 / Cost, cleanup and interview story

Understand why a safe learning environment still needs ownership.

Cost model

The selected ceiling is USD10/month, not a cap. A customer-managed key has a base monthly charge; the generated secret, storage and API usage add costs. Credits affect amount due, not whether usage happened. Optional Config and GuardDuty need their own usage estimate and exercise window.

Retention tradeoff

The template retains both buckets, the key and secret on stack deletion. Bucket lifecycle expires learning data after the documented windows. Retention protects against accidental stack-driven data loss but can preserve cost; key deletion can make remaining evidence unreadable. Review object versions, delete markers and dependency/recovery windows before cleanup.

A 60-second answer

"I deployed a Sydney AWS baseline with private SSE-KMS buckets, a rotating key, management-event CloudTrail, two Config rules and a demonstration secret. Nineteen configuration checks and six behavior checks passed. I matched real events in delivered logs and closed an account S3 guardrail gap. I validated one digest and 17 log files. Locally, I upgraded an nginx target and repeated the same OpenVAS checks, reducing candidate alarms from two to zero. I also tested a synthetic GoPhish workflow. MFA and restricted human access remain unverified."

A strong troubleshooting answer

Describe the symptom, narrow the cause, make the smallest relevant fix and show the comparable retest. The SQLite cleanup or SMTP From validation issues are useful engineering examples. S3-001 is the clearer security-remediation example.

What an interviewer may challenge

Ask: Which identity was tested? Which Region? Which resource population? Which event or finding? What time range? What result proves closure? What remains unverified? Answer those before making broader claims about compliance or account-wide hardening.

09 / Hands-on command reference

Run from the project root with an approved identity. Replace placeholders; never paste credentials.

1. Check identity and code

In CloudShell, the existing console identity is already available. On your computer use the approved short-lived profile described in the deployment guide. Confirm the account before continuing.

```
aws sts get-caller-identity
```

```
cfn-lint infra/baseline.json infra/iam-auditor.json
```

```
python -m unittest discover -s tests -v
```

```
aws cloudformation validate-template --template-body file://infra/baseline.json --region
ap-southeast-2
```

2. Prepare, review and deploy

The first command prepares a new stack only. Read describe-change-set before execute-change-set. For an existing stack use UPDATE with reviewed parameters; do not rerun a CREATE workflow blindly.

```
aws cloudformation create-change-set --stack-name aws-security-lab --change-set-name baseline-review
--change-set-type CREATE --template-body file://infra/baseline.json --parameters
file://infra/parameters-project.json --capabilities CAPABILITY_IAM --region ap-southeast-2
```

```
aws cloudformation describe-change-set --stack-name aws-security-lab --change-set-name
baseline-review --region ap-southeast-2
```

```
aws cloudformation execute-change-set --stack-name aws-security-lab --change-set-name
baseline-review --region ap-southeast-2
```

3. Retest and retain the result

Use a new output name each time. The behavior check writes harmless lab data and reads the demo value into memory. The local simulation requires Docker and sends no real email.

```
python scripts/verify_live.py --stack aws-security-lab --region ap-southeast-2 --expected-account
YOUR_ACCOUNT --output private-evidence/live-01.json
```

```
python scripts/verify_behavior.py --stack aws-security-lab --region ap-southeast-2
--expected-account YOUR_ACCOUNT --output private-evidence/behavior-01.json
```

```
python scripts/run_local_simulation.py
```

Question bank / 01

Questions 1-4 | Answer aloud, then compare with the outline.

01. What problem does the lab solve?

It makes security settings repeatable and gives each claim a verification step. It does not turn an account into a certified organization.

02. Why CloudFormation?

One declared stack shows what will be created. A change set makes additions and permission changes reviewable before execution.

03. How are IAM and MFA different?

IAM policies authorize actions/resources. MFA strengthens a sign-in or role-assumption path. A highly privileged session can still use MFA.

04. How does the auditor enforce least privilege?

It can list/read one exercise bucket and decrypt those objects through S3. It cannot write, read the log bucket or retrieve the secret. Effective permissions need live testing.

Question bank / 02

Questions 5-8 | Answer aloud, then compare with the outline.

05. Is the auditor your current console identity?

No. It is an optional IAM-user assumption exercise for a suitable environment. The current project uses Builder ID and an AWS-managed human role.

06. Why not apply an MFA condition everywhere?

Federation and IAM-user context differ. MultiFactorAuthPresent is not a universal indication of upstream MFA. Test the actual identity path.

07. What does CloudTrail record?

This template selects read/write management events. It does not select S3 object data events or every possible event category.

08. Does a multi-Region trail make every service multi-Region?

No. Config and GuardDuty are configured and verified regionally. Even a multi-Region trail requires regional delivery/coverage evidence.

Question bank / 03

Questions 9-12 | Answer aloud, then compare with the outline.

09. Why enable log validation?

CloudTrail delivers signed digests for integrity checking. This lab validated 1/1 digest and 17/17 log files in the recorded interval. Earlier attempts found no digests in the interval and stayed incomplete; enabling the setting alone was not enough.

10. What does Config add?

Resource configuration history and evaluations against selected rules. It is detective, does not prevent changes, and has explicit resource/Region scope.

11. What if Config returns no results?

The test is incomplete. INSUFFICIENT_DATA, stale evaluations or an empty report do not mean compliant.

12. What does GuardDuty add?

Managed threat detections from supported telemetry. A sample finding tests handling; it is not proof that a real attack was detected.

Question bank / 04

Questions 13-16 | Answer aloud, then compare with the outline.

13. Why both public-access controls and encryption?

They solve different problems. An encrypted object can still be disclosed to an authorized or overly broad caller. Encryption does not fix a public policy.

14. Why a customer-managed KMS key?

It makes key policy, rotation and permissions inspectable. It adds cost and a separate availability risk if disabled/deleted.

15. Why do the service grants have SourceArn/SourceAccount conditions?

They constrain service use to the intended trail/account/Region and reduce confused-deputy exposure. They must still match the actual service request context.

16. Does default encryption rewrite old data?

No. Existing objects and versions need separate review; the upload test checks new-object behavior.

Question bank / 05

Questions 17-20 | Answer aloud, then compare with the outline.

17. Why retain the key and buckets on stack deletion?

To avoid accidental evidence loss or unreadable retained data. Retention also leaves costs and a separate cleanup obligation.

18. Why Secrets Manager instead of a configuration file?

It centralizes secret access and auditing. The demonstration uses a generated value. Rotation and retrieval boundaries require additional tests.

19. Does the budget enforce $10?

No. It tracks cost and sends alerts when operational; billing can lag. In this lab the recipient was inactive, so alert delivery remains an open gap.

20. How do you prove vulnerability closure?

Preserve the original report, confirm the condition, remediate it and repeat comparable checks. Here nginx 1.27.5 became 1.30.6; both runs detected the service, and the same targeted selection went from two low-confidence alarms to zero. NGINX-001 closes the old-version condition, not every possible host vulnerability.

Question bank / 06

Questions 21-24 | Answer aloud, then compare with the outline.

21. What does the local phishing test prove?

Campaign creation, SMTP capture and click-event recording work for a synthetic recipient. An automated click says nothing about human susceptibility.

22. What is SOC 2 alignment here?

The controls can support access, monitoring and change/configuration objectives, but a small lab cannot establish organizational SOC 2 compliance or certification.

23. What failed during the build?

Budget contact activation was denied and remains open. Local campaign network/SQLite/From-address issues were fixed and retested. Config initially stalled on recorder/channel startup order; rollback, explicit stopped creation and later start resolved it. First delivery/evaluation checks also ran too early; later fresh checks passed.

24. What would you improve for production?

Organizational identity/MFA enforcement, regional coverage, durable retention, alert routing, incident ownership, backups and regularly reviewed effective permissions. Scope and price these explicitly.

Question bank / 07

Questions 25-28 | Answer aloud, then compare with the outline.

25. An API says AccessDenied. Is the control working?

Not necessarily. Identify the caller and action, and check identity policy, trust, resource policy, session policy, boundary and organization controls. Pair the denied operation with the intended allowed path to isolate the cause.

26. The trail is logging but there are no delivered objects. What next?

Inspect status errors, bucket prefix/policy, KMS key state/policy and delivery timing. Generate a benign management event and correlate it after delivery. Keep the operational test incomplete until records arrive.

27. Why can a green Config dashboard be misleading?

The resource may be out of scope, the recorder stopped, evaluations stale, or results absent. Check resource population, time, delivery and rule evaluation details.

28. A secret retrieval test passed as administrator. Is that least privilege?

No. It demonstrates only that the administrator can retrieve it. Test a deliberately scoped authorized identity and an appropriate restricted identity, without printing the value.

Question bank / 08

Questions 29-32 | Answer aloud, then compare with the outline.

29. What if the scanner finds nothing?

Check host reachability, alive detection, feed state, authentication, ports and scan completion before interpreting results. An empty output may mean inadequate coverage.

30. How do you choose a finding severity?

Use verified exposure, affected asset, exploit conditions and impact. Preserve the scanner score separately from contextual risk. Do not invent a severity for an unconfirmed issue.

31. Can you claim phishing success from opens?

Opens and clicks can be generated by automation and mail-security systems. This lab explicitly used an automated click, so it measures workflow operation rather than human response.

32. Why separate logs from exercise evidence?

It makes delivery policies and test access easier to reason about. The auditor exercise can read only harmless exercise data rather than grant access to all activity logs.

Question bank / 09

Questions 33-36 | Answer aloud, then compare with the outline.

33. Would you use the 30-day lifecycle in production?

Only if it matched an approved retention requirement. Audit, legal and incident needs can require different retention, deletion controls and immutability. The lab uses a learning-specific window.

34. What is a CloudFormation change set?

A proposed set of resource changes. Review additions, replacements, deletions, parameters and IAM effects before execution. Ready-to-execute status is different from deployed stack success.

35. What would you do after a stack rollback?

Read the failure events, identify the failing service/policy and inventory retained or partially created resources. Do not assume rollback removed all cost or evidence.

36. Does Block Public Access remove every S3 risk?

No. It addresses specific public-policy/ACL paths. Overly broad authorized access, cross-account design, data handling, key access and other configuration still need review.

Question bank / 10

Questions 37-40 | Answer aloud, then compare with the outline.

37. How do you avoid exposing evidence in GitHub?

Keep raw outputs, credentials, screenshots and databases in ignored private storage. Publish sanitized summaries with scope, UTC dates, methods and limitations. Review the packaged files before publication.

38. What is the difference between implemented and effective?

Implemented means the configuration exists. Effective means the relevant behavior worked for the intended population and period. The tests and evidence need to match that claim.

39. What should you say when a feature is unsupported?

State the exact environment limitation, its effect on the claimed control and the supported setup required. Do not mark an unavailable service as enabled or convert the account without the owner decision.

40. What remains before the original resume entry is fully defensible?

Live deployment and scoped Config results are verified. The owner excluded GuardDuty and multi-Region operation in this account, so remove those claims. MFA and restricted-human-role enforcement remain unverified. The completed targeted nginx scan and retest support version-condition closure, not a full-host assessment or two proven exploits. Disclose inactive budget alerts.

Practice / Evidence before claims

A short preparation routine and official references.

20-minute mock interview

2 minutes: introduce the project and boundaries. 5 minutes: draw the architecture. 5 minutes: explain S3-001 from observation to persisted retest. 5 minutes: answer IAM/logging failure scenarios. 3 minutes: explain cost, cleanup and remaining gaps.

Score each answer

Give yourself 0-2 for technical accuracy, actual project evidence, reasoning and clarity. A strong answer names the tested scope and limit. A weak answer repeats a definition without explaining a result.

Before the interview

Open the current evidence index, findings register and test report. Use the latest status, not an old screenshot or planned procedure. Do not claim production hardening, real-person awareness results or vulnerability closure that was not observed.

AWS account experience and human-access model

Supported services and regional restrictions

Advanced activation prerequisites and consequences

CloudTrail KMS permissions

CloudTrail integrity validation

Greenbone Community Containers

GoPhish campaign API